

CIIC 4050 / ICOM 5007 Operating Systems

Lecture 4: Interprocess Communication

Tanenbaum & Woodhull, Operating Systems: Design and Implementation, © 2006, Prentice-Hall, Inc. All rights reserved. 0-13-142938-8

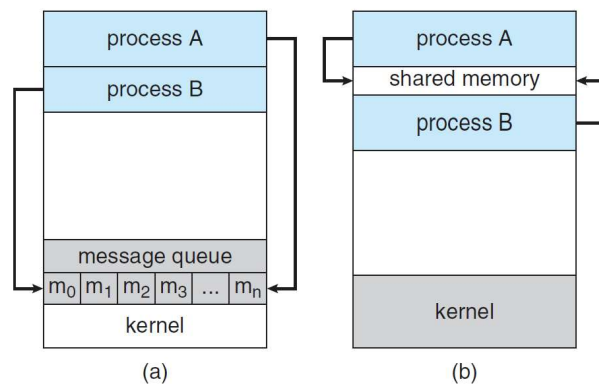
41

Interprocess Communication (IPC)

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
- Cooperating processes need **IPC**
- Two models:
 - Shared memory
 - Message passing

42

Communications Models



43

Cooperating Processes

- Independent** process cannot affect or be affected by the execution of another process
- Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

44

Producer-Consumer Problem (1)

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - unbounded-buffer** places no practical limit on the size of the buffer
 - bounded-buffer** assumes that there is a fixed buffer size

45

Producer-Consumer Problem (2)

- Shared data


```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```
- Can only use $\text{BUFFER_SIZE} - 1$ elements

46

Bounded Buffer - Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out);
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

NOTE: This solution might suffer from synchronization problems that we shall address later in the course ... but for the moment ...

47

Bounded Buffer - Consumer

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```

48

IPC – Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details in Chapter 5.

49

IPC – Message Passing (1)

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(message) – message size fixed or variable
 - **receive**(message)

50

IPC – Message Passing (2)

- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

51

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

52

Direct Communication

- Processes must name each other explicitly:
 - **send** (P , *message*) – send message to process P
 - **receive**(Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

53

Indirect Communication

- Messages using mailboxes (or ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

54

Indirect Communication

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - **send**(A , *message*) – send a message to mailbox A
 - **receive**(A , *message*) – receive a message from mailbox A

55

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

56

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **send** sender blocks until the message is received
 - **receive** receiver blocks until a message is available
- **Non-blocking** is considered **asynchronous**
 - **send** sender sends message and continue
 - **receive** receive a valid message or null

57

Buffering

- Queue of messages attached to the link; implemented in one of three ways
 - Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 - Bounded capacity – finite length of n messages
Sender must wait if link full
 - Unbounded capacity – infinite length
Sender never waits

58

Examples of IPC Systems - POSIX

- Process first creates shared memory segment
 - `segment id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);`
 - See also `shm_open()`
- Process wanting access to that shared memory must attach to it
 - `shared memory = (char *) shmat(id, NULL, 0);`
- Now the process could write to the shared memory
 - `sprintf(shared memory, "Writing to shared memory");`
- When done a process can detach the shared memory from its address space
 - `shmdt(shared memory);`

59

IPC Systems – Windows XP (1)

- Message-passing centric via **local procedure call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels

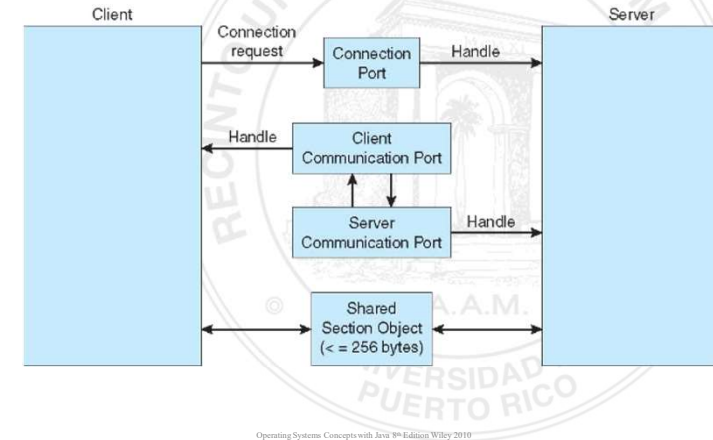
60

IPC Systems – Windows XP (2)

- Communication works as follows:
 - The client opens a handle to the subsystem's connection port object
 - The client sends a connection request
 - The server creates two private communication ports and returns the handle to one of them to the client
 - The client and server use the corresponding port handle to send messages or callbacks and to listen for replies

61

LPCs in Windows XP



62

Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

63

Sockets

- A socket is defined as an *endpoint for communication*
- Concatenation of IP address and port
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a (unique) pair of sockets

64

Sockets (2)

- Two widely used address domains:
 - unix domain**, 2 processes share a common file system
 - internet domain**, 2 processes running on any two hosts on the internet communicate
- Two widely used socket types
 - stream sockets**, treat communications as a continuous stream of characters
 - datagram sockets** read entire messages at once

65

Socket Communication (Server)

```
/* A simple server using TCP Sockets
   The port number used is passed as an argument */
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main(int argc, char *argv[])
{
    int sockfd, newsockfd, portno, clien;
    char buffer[256];
    struct sockaddr_in serv_addr, cli_addr;
    int n;
    char *z;

    FILE *file;
    if (argc < 2) {
        fprintf(stderr, "ERROR, no port provided\n");
        exit(1);
    }

    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0)
        error("ERROR opening socket");

    bzero((char *) &serv_addr, sizeof(serv_addr));
    portno = atoi(argv[1]);
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = INADDR_ANY;
    serv_addr.sin_port = htons(portno);
```

66

Socket Communication (Server)

```
if (bind(sockfd, (struct sockaddr *) &serv_addr,
    sizeof(serv_addr)) < 0)
    error("ERROR on binding");

listen(sockfd,5);
clilen = sizeof(cli_addr);

while((newsockfd = accept(sockfd,(struct sockaddr *) &cli_addr,
    &clilen)) >= 0){
    if (newsockfd < 0)
        error("ERROR on accept");

    bzero(buffer,256);

    n = read(newsockfd,buffer,255);
    if (n < 0) error("ERROR reading from socket");

    printf("Here is the message: %s\n",buffer);
    z = inet_ntoa(*(struct in_addr *)&cli_addr); /* cast as a struct in_addr */

    printf("Client information: %s\n", z);

    file = fopen("clients.txt","a+");
    fprintf(file,"%s\t%s",z,buffer);
    fclose(file);

    n = write(newsockfd,"I got your message\nMessage: %",18);
    if (n < 0) error("ERROR writing to socket");
}
printf("Server Terminated\n");
return 0; }
```

67

Socket Communication (Client)

```
/* A simple client using TCP Sockets
   The server address and port number
   are passed as arguments */
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

int main(int argc, char *argv[])
{
    int sockfd, portno, n;
    struct sockaddr_in serv_addr;
    struct hostent *server;

    char buffer[256];
    if (argc < 3) {
        fprintf(stderr, "usage %s hostname port\n", argv[0]);
        exit(0);
    }

    portno = atoi(argv[2]);
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0)
        error("ERROR opening socket");

    server = gethostbyname(argv[1]);
    if (server == NULL) {
        fprintf(stderr, "ERROR, no such host\n");
        exit(0);
    }
}
```

68

Socket Communication (Client)

http://www.linuxhowtos.org/C_C++/socket.htm

```

bzero((char *) &serv_addr, sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
bcopy((char *)server->h_addr,
      (char *)&serv_addr.sin_addr.s_addr,
      server->h_length);
serv_addr.sin_port = htons(portno);

if (connect(sockfd,&serv_addr,sizeof(serv_addr)) < 0)
    error("ERROR connecting");
printf("Please enter the message: ");

bzero(buffer,256);
fgets(buffer,255,stdin);
n = write(sockfd,buffer,strlen(buffer));
if (n < 0)
    error("ERROR writing to socket");

bzero(buffer,256);
n = read(sockfd,buffer,255);
if (n < 0)
    error("ERROR reading from socket");
printf("%s\n",buffer);
return 0;
}

```

Tanenbaum & Woodhull, Operating Systems: Design and Implementation, © 2006 Prentice-Hall, Inc. All rights reserved. 0-13-142938-8

69

Socket Communication in Java (1)

```

public class DateServer {
    public static void main (String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            // now listen for connections
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                // write the Date to the socket
                pout.println(new java.util.Date().toString());

                // close the socket and resume
                // listening for connections
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}

```

Operating Systems Concepts with Java 8th Edition Wiley 2010

70

Socket Communication in Java (2)

```

public class DateClient {
    public static void main (String[] args) {
        try {
            // make connection to server socket
            Socket sock = new Socket("127.0.0.1", 6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            // read the Date from the socket
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            // close the socket connection
            sock.close();
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}

```

Operating Systems Concepts with Java 8th Edition Wiley 2010

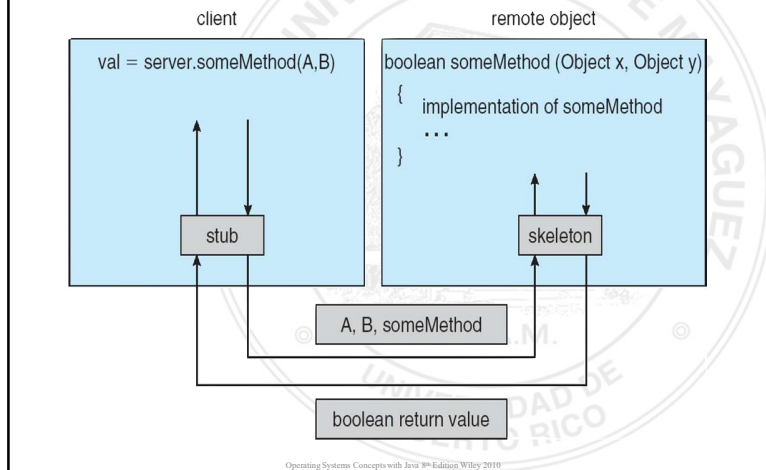
71

Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems; built on top of local IPC mechanism
- **Stubs** – client-side proxy for the actual procedure on the server, typically one stub per each remote procedure
- The client-side stub locates the server and *marshalls* the parameters (packaging in a form that can be transmitted, endianness...)
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server

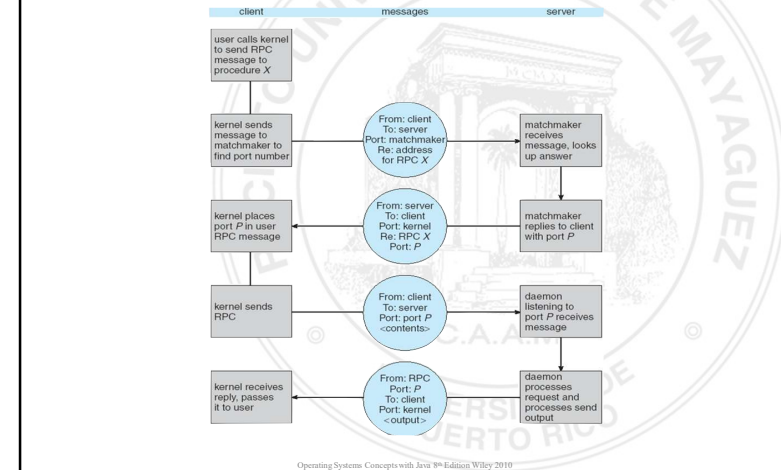
72

Marshaling Parameters




73

Execution of RPC



74

Remote Method Invocation



- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs
- RMI allows a Java program on one machine to invoke a method on a remote object

75

RMI Example

- Begin by declaring an interface that specifies the methods that can be invoked remotely...

```
public interface RemoteDate extends Remote {
    public abstract Date getDate() throws RemoteException;
}
```

76

RMI Example

```
public class RemoteDateImpl extends UnicastRemoteObject
    implements RemoteDate {
    public RemoteDateImpl() throws Remote Exception {}

    public Date getDate() throws Remote Exception {
        return new Date();
    }

    public static void main(String[] args) {
        try {
            Remote Date dateServer = new RemoteDateImpl();

            // Bind this object instance to the name "DateServer"
            Naming.rebind("DateServer", dateServer);
        } catch (Exception e) {
            System.err.println(e);
        }
    }
}
```

Operating Systems Concepts with Java 8th Edition Wiley 2010

77

RMI Example

```
public class RMIClient {
    public static void main(String[] args) {
        try {
            String host = "rmi://127.0.0.1/DateServer";

            RemoteDate dateServer = (RemoteDate) Naming.lookup(host);
            System.out.println(dateServer.getDate());
        } catch (Exception e) {
            System.err.println(e);
        }
    }
}
```

Operating Systems Concepts with Java 8th Edition Wiley 2010

78