

## CIIC 4050 / ICOM 5007 Operating Systems

### Lecture 5: Threads

Tanenbaum & Woodhull, Operating Systems: Design and Implementation, © 2006, Prentice-Hall, Inc. All rights reserved. 0-13-142938-8

79

## Threads (1)

- A **thread** is an execution context for a processor, which handles an execution of the instructions in a process or part of them.
- In a system that does not support multiple threads in a process, each process has only one thread.
- In a system that supports multiple threads in a process, there can be many threads executing concurrently on the same process.

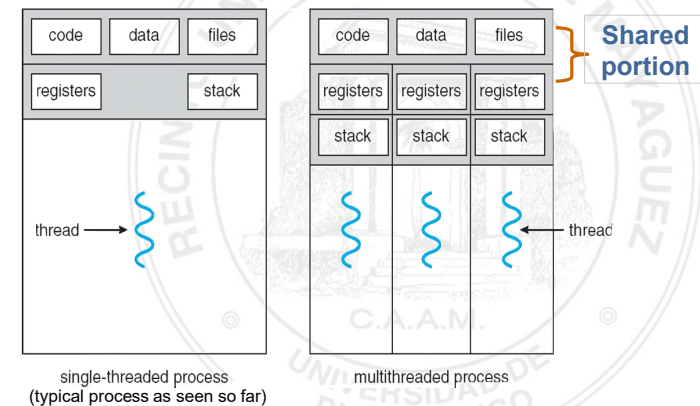
80

## Threads (2)

- Each thread in a process
  - can be executing a different portion (method or subroutine) of the process's code.
- In a uniprocessor system,
  - they alternate execution as different processes do.
- In a multiprocessors system, they can execute in parallel.
- For each thread, the following data is maintained:
  - *ID*, *program counter*, *register set*, and *stack*. This is in addition to the other information maintained in PCB of the owner process.

81

## Single and Multithreaded Processes



82

## Thread Content

Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

- **Figure 2-7.** The first column lists some items shared by all threads in a process. The second one lists some items private to each thread.

Tanenbaum & Woodhull, Operating Systems: Design and Implementation, © 2006, Prentice-Hall, Inc. All rights reserved. 0-13-142938-8

83


## Thread use

- Potential uses of threads:
  - Word Processor – 1 thread display, 1 thread keystrokes, 1 thread spelling & grammar
  - Web browser – 1 thread display images, 1 thread network
  - Web server – ???

Tanenbaum & Woodhull, Operating Systems: Design and Implementation, © 2006, Prentice-Hall, Inc. All rights reserved. 0-13-142938-8

84

## Example: Multithreaded Server Architecture



### General Algorithm followed by server:

1. Initialization
2. While (true)
  - a. get request
  - b. initiate thread to serve request concurrently

85

## Benefits (1)

- **Responsiveness** - threads allow a program to continue executing even if part of it is blocked or is performing a lengthy computation. This may result in faster responses during program execution.
- **Resource Sharing** - data and resources can be shared in an efficient manner. That kind of sharing in processes is achieved by message passing or by shared memory, and at a major cost in terms of effort from the programmer.

86

## Benefits (2)

- **Economy** - several of the overhead issues involved in creating and managing processes are highly minimized in threads.
- **Scalability** - benefits of multithreading can be greatly increased in multiprocessor architectures.

87

## Benefits (3)

- **Thread vs Process Cost\***

	User threads	LWP/kernel threads	Processes
Creation time	52	350	1700
Synchronization with semaphores	66	390	200

\* on SPARC station 2 (Solaris), from Unix Internals by Uresh Vahalia, PH 1996.

88

## Multicore Programming (1)

- Multicore or multiprocessor systems putting pressure on OS designers as well as on application programmers.
- Challengers for OS designers include:
  - Scheduling to allow efficient use of cores and parallel execution
  - Synchronization of threads running in different cores.


89

## Multicore Programming (2)

- Challengers for application programmers include:
  - **Dividing activities** - which can be parallelized
  - **Balance** - equal effort on each parallel activity
  - **Data splitting** - divide data managed by each activity
  - **Data dependency** - between different activities and properly synchronize responsible threads.
  - **Testing and debugging** - involves more effort than in sequential processes since order of execution of parallel tasks may be undetermined.

90


## Concurrent Execution on a Single-core System



- Concurrency = Interleaved execution (time sharing systems)

91

## Parallel Execution on a Multicore System



- Concurrency = Parallel execution (one thread/core + time sharing)

92

## Management of Threads (1)

- Thread management can be provided at either:
  - *user level* - are supported above the kernel
  - *kernel level* - are supported and managed directly by the OS.
- Virtually all modern OS's support kernel threads
  - Windows XP/2003, Solaris, Linux, Tru64 UNIX, HP-UX, Mac OS X

93

## Management of Threads (2)


- Ultimately, there must be a relationship between a running user thread and a kernel thread to allow the execution of that user thread
- Different models are used to match user threads to corresponding kernel threads for their execution

94

## Many-to-One

- Many user-level threads mapped to single kernel thread
- Examples:
  - Solaris Green Threads
  - GNU Portable Threads

(-) Threads execute concurrently, but not in parallel...  
 (-) Threads are scheduled by process itself.




95

## One-to-One

- Each user-level thread maps to kernel thread
- Examples
  - Windows NT/XP/2000
  - Linux
  - Solaris 9 and later

Different threads in a process can run in parallel. The kernel schedules them...


(-) more overhead since kernel needs to create one kernel thread for each user thread created



96

## Many-to-Many Model


- Allows many user level threads to be mapped to many kernel threads (multiplexed)
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the ThreadFiber package



97

## Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
  - IRIX
  - HP-UX
  - Tru64 UNIX
  - Solaris 8 and earlier



98



## Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS
- Example of thread libraries:
  - POSIX **Pthreads**
  - Win32 threads
  - Java threads

99

## PTHREADS

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library (specification), implementation is up to development of the library
- Commonly implemented in UNIX operating systems (Solaris, Linux, Mac OS X)

100

## PTHREADS (example)

```
#include <pthread.h>
#include <stdio.h>
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */
int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of attributes for the thread */
    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        /*exit(1);*/
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "Argument %d must be non-negative\n", atoi(argv[1]));
        /*exit(1);*/
        return -1;
    }
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
```

101

## PTHREADS (example)

```
/* now wait for the thread to exit */
pthread_join(tid, NULL);
printf("sum = %d\n", sum);
}
/**
 * The thread will begin control in this function
 */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;
    if (upper > 0) {
        for (i = 1; i <= upper; i++)
            sum += i;
    }
    pthread_exit(0);
}
```

102

## Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threading model provided by underlying OS
- Java threads may be created by:
  - Extending `Thread` class
  - Implementing the `Runnable` interface (most common)

103

## Java Threads (1)

- Runnable Interface

```
public interface Runnable {
    void run();
}
```

- A class implementing the interface defines objects corresponding to threads
  - once instantiated, the thread object will start executing its `run()` method...

104

## Java Threads (2)

- To create and start the execution of a thread object of type `Runnable`, the following is required:
  1. Create an object of type `Thread` while passing an instance of the `Runnable` object as parameter of constructor.
  2. To that object of type `Thread` apply method `start()`.

\*\*\* The new thread will automatically start executing the `run()` method of the associated `Runnable` object.

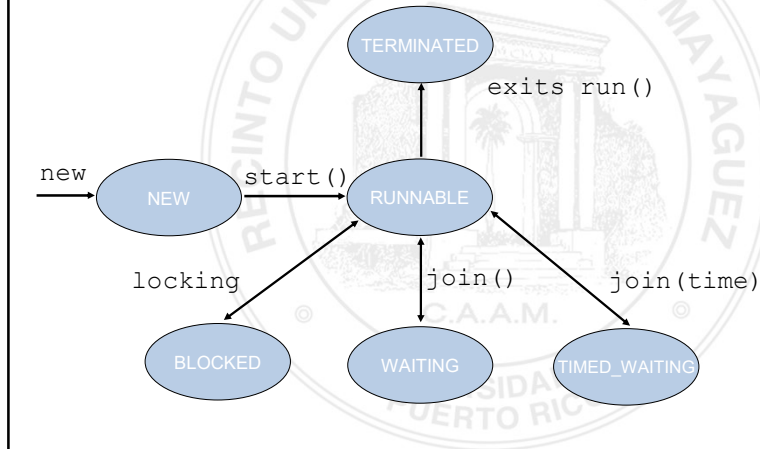
105

## Java Threads States

- A Java thread may be in one of **6 possible states**
  - **New** - Thread created but `run()` method has not started
  - **Runnable** - when the `run()` method starts
  - **Blocked** - thread is waiting for a lock
  - **Waiting** - thread is waiting for an action of another thread (example: executing `join()`)
  - **Timed waiting** - same as waiting for for a specified maximum time. Useful to eliminate the possibility of starvation...
  - **Terminated** - when the run method of the thread has finished.

106

## Java Threads States



107

## Threads: Sum of Array Elements (1)

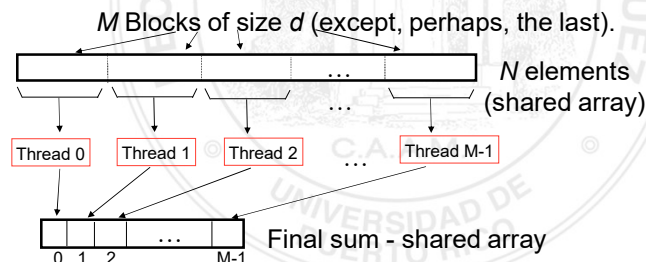
- Given an array of  $N$  integers and a fixed value  $d$ , create
- $M$  threads to add all its elements in portions of size  $d$ .
- One possible solution:
  - Main thread does the following
    - Create several thread
    - Each thread is assigned a portion of the array
    - Each thread computes the sum of its assigned portion
    - Results are placed in appropriate entries in a shared array

108

## Threads: Sum of Array Elements (2)

- Given an array of  $N$  integers and a fixed value  $d$ , create  $M$  threads to add all its elements in portions of size  $d$ .

$$M = \begin{cases} N/d & \text{if } N \% d == 0 \\ N/d + 1 & \text{if not} \end{cases}$$



109

## Java: Sum of Array Elements (1)

```

public static void main(String[] args) {
    int[] a = ...; // the array to add
    int pSize = ...; // size of blocks
    int nThreads =
        (a.length % pSize == 0 ? a.length / pSize : a.length / pSize + 1);

    Barrier barrier = new Barrier(nThreads);

    int[] sum = new int[nThreads];

    for (int i=0; i<nThreads; i++) {
        Thread t =
            new Thread(new ArrayPortionAdder(a, sum, i, pSize, barrier));
        t.start();
    }

    // wait for all threads to finished as per the barrier object...
    try {
        barrier.waitForRelease();
    }
    catch (InterruptedException e) {}

    // compute the final sum from the array sum ...
    ...
}
  
```

110



## Java: Sum of Array Elements (2)

```
public class ArrayPortionAdder implements Runnable {
    private int[] arr, sum;
    private int index, size;
    private Barrier barrier;

    public ArrayPortionAdder(int[] arr, int[] sum, int index, int size,
        Barrier b) {
        ... arr is the array to sum, sum is the array where partial results are placed, ...
    }

    public void run() {
        int low = index * size;
        int sup = low + size - 1;
        if (sup >= arr.length)
            sup = arr.length - 1;
        int asum = 0;

        for (int i = low; i <= sup; i++) // thread sums its part
            asum += arr[i];

        sum[index] = asum; // places result in shared array
        // count one more thread as finished
        ...
        barrier.incCount();
        ...
    }
}
```

111

## Java: Object Type for Synchronization

```
public class Barrier {
    private int threshold, count = 0;
    public Barrier(int t) {
        threshold = t;
    }
    public void reset() {
        count = 0;
    }
    public synchronized void waitForRelease() throws InterruptedException {
        while (count < threshold)
            wait();
    }
    public synchronized void incCount()
        throws InterruptedException {
        count++;
        if (count == threshold)
            notifyAll();
    }
}
```

\*\*\* This is an example of a class used for synchronization – we shall study this in more detail later in the course....

112

## Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Thread cancellation of target thread
  - Asynchronous or deferred
- Signal handling
- Thread pools
- Thread-specific data
- Scheduler activations

113

## Semantics of fork() and exec()

- Does fork() duplicate only the calling thread or all threads?
- Some UNIX systems have two versions of fork():
  - one that duplicates all threads active in the parent process
  - one that duplicates only the thread invoking the fork()

114

## Thread Cancellation (1)

- Terminating a thread before it has finished
  - Multiple threads performing database search, one returns value, others might be canceled...
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
    - Case where resources have been allocated to canceled thread, thread canceled while updating data needed by other threads
    - java – stop() (deprecated)

115

## Thread Cancellation (2)

- **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
  - pthreads – cancellation points
  - java – interrupt(), isInterrupted()

116

## Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred
- A **signal handler** is used to process signals.
- Signal handling follows the pattern:
  - Signal is generated by particular event
  - Signal is delivered to a process
  - Signal is handled

117

## Signal Handling (example)

```
/* Sample program to handle system signals */
#include <signal.h>
#include <stdio.h>

void * myhandler(int myint)
{
    printf("\nSignal Handled!!\n\n");
    exit(0);
}

int main()
{
    signal( SIGINT, (void *) myhandler );
    signal( SIGTERM, (void *) myhandler );

    while(1) {
        printf("Doing Nothing...\n");
        sleep(1);
    }
}
```

118

## Signal Handling

- Options:
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals for the process

119

## Thread Pools (1)

- To create a new thread, although simpler than to create a process, has its impacts on the system:
  - To create a new thread consumes time
  - Too many active threads may negatively impact system's performance

120

## Thread Pools (2)

- An alternative is to create a number of threads in a pool where they await work. This has advantages:
  - Usually slightly faster to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool

121

## Thread Pools in Java (1)

- Java allows the creation of thread pools as shown next.
- Java classes to support thread pools:

```
public interface Executor {
    // executes the given thread
    // at some time in the future
    void execute(Runnable t);
}
```

122

## Thread Pools in Java (2)

- Class **Executors** - provides a set of useful methods to create objects that are relevant to thread management in JVM. One static method creates a pool of a given number of threads. Such pool is created as an object of type **ExecutorService**.

```
public static ExecutorService
    newFixedThreadPool(int maxNoThreads)

Interface ExecutorService
    public interface ExecutorService extends Executor {
        ... several methods + execute ...
    }
```

123

## Thread Pools in Java (2)

- Example of a Java thread pool:
- Assume that class Worker is as follows:

```
public class Worker implements Runnable {
    private static int NT = 0;
    public void run() {
        System.out.println("Thread number " + (++NT));
    }
}
```

- The following creates a thread pool:

```
Runnable r1 = new Worker();
Runnable r2 = new Worker();
Runnable r3 = new Worker();
ExecutorService pool = Executors.newFixedThreadPool(3);
pool.execute(r1);
pool.execute(r2);
pool.execute(r3);
```

124

## Thread Specific Data


- Allows each thread to have its own copy of data
  - Transaction processing system (each thread handle separate transaction)
- Useful when you do not have control over the thread creation process (i.e., when using thread pools)

Transaction	Thread
new_order()	trans_T1
payment()	trans_T3
order_status()	trans_T5
delivery()	trans_T2
stock_value()	trans_T4

125

## Scheduler Activations

- Both M:M and Two-level models require communication to dynamically maintain the appropriate number of kernel threads allocated to the application
- Place a data structure between user and kernel threads (lightweight process)
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
- This communication allows an application to maintain the correct number kernel threads



126

## Windows XP Threads (1)

- Implements the one-to-one mapping, kernel-level
- Each thread contains
  - A thread id
  - Register set
  - Separate user and kernel stacks
  - Private data storage area

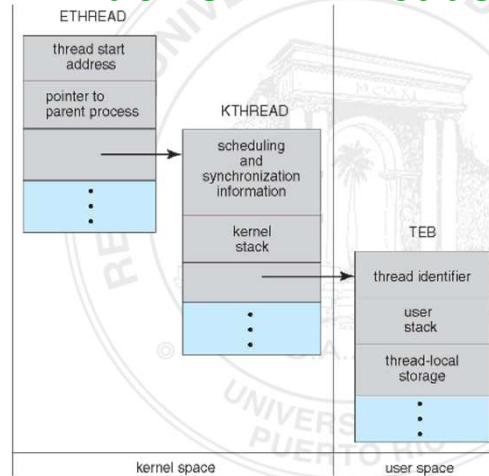
127

## Windows XP Threads (2)

- The register set, stacks, and private storage area are known as the **context** of the threads
- The primary data structures of a thread include:
  - ETHREAD (executive thread block)
  - KTHREAD (kernel thread block)
  - TEB (thread environment block)

128

## Windows XP Threads



129

## Linux Threads (1)

- Linux refers to them as *tasks* rather than *threads* or *processes*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process), several flags control amount of sharing between parent-child, if no flags are set:
- **clone() = fork()**

130



## Linux Threads (2)

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

131